

PGCF: Physiology-Guided Consistency Framework for Deepfake Detection via Multi-Stream Temporal Attention

VSS Pavan Ganesh
Dept. of Computer Science & Engineering
IIIT Naya Raipur
Naya Raipur, India
vss24100@iiitnr.edu.in

Suraj Meruva
Dept. of Data Science & AI
IIIT Naya Raipur
Naya Raipur, India
meruva24102@iiitnr.edu.in

Abhiram Boini
Dept. of Data Science & AI
IIIT Naya Raipur
Naya Raipur, India
boini24102@iiitnr.edu.in

Abstract—The proliferation of deepfake media poses a significant threat to digital trust, demanding robust and interpretable detection methods. We present the Physiology-Guided Consistency Framework (PGCF), a three-stream neural architecture that fuses spatial, frequency-domain, and remote photoplethysmography (rPPG) signals for deepfake video detection. The key contribution is the Physiology-Guided Cross-Modal Consistency (PGCMC) module, which gates multi-stream feature fusion using rPPG-derived spectral entropy as a physiological confidence signal. The spatial stream employs a ResNet50 backbone augmented with a Convolutional Block Attention Module (CBAM) for spatially-aware artifact localization, while temporal modeling is enhanced by a Multi-Head Self-Attention layer inserted before a Bidirectional LSTM to capture long-range frame dependencies. Evaluated on a balanced dataset of 356 video clips (178 real, 178 fake), PGCF achieves an AUC of 0.8240, accuracy of 74.16%, and Average Precision of 0.7907. Analysis of cross-modal consistency scores reveals a strong discriminative gap: real videos attain a mean consistency of 0.690 versus 0.272 for fake videos. Robustness evaluation under Gaussian noise (σ up to 0.20) demonstrates graceful AUC degradation to 0.73, validating model resilience to input perturbations.

Index Terms—deepfake detection, remote photoplethysmography, multi-stream architecture, cross-modal consistency, CBAM, BiLSTM, temporal self-attention, face forgery

I. INTRODUCTION

Face manipulation technologies, colloquially termed *deepfakes*, have reached a level of visual fidelity that renders pixel-level inspection by human observers unreliable. Generative Adversarial Networks (GANs), autoencoder-based face-swapping, and neural rendering pipelines can produce convincing videos, with direct implications for misinformation, identity fraud, and political manipulation [1].

Existing detection approaches fall broadly into three categories: (i) *spatial artifact* methods that exploit GAN fingerprints and blending inconsistencies in the RGB domain [2]; (ii) *frequency-domain* methods that detect DCT or Fourier abnormalities introduced by generation [3]; and (iii) *physiological signal* methods that exploit the absence of coherent biological cues (e.g., heartbeat-induced skin color variation)

in synthesized faces [5]. Each paradigm has blind spots: RGB methods are sensitive to compression; frequency methods fail on high-bitrate re-encoding; physiological methods degrade under motion or illumination change.

We address this limitation with the **Physiology-Guided Consistency Framework (PGCF)**, which fuses all three streams under a physiological gating mechanism. Genuine faces exhibit strong cross-modal consistency—spatial textures, frequency patterns, and rPPG signals align—whereas synthesized faces break this coherence in at least one modality. The PGCMC module quantifies this alignment and uses rPPG spectral confidence to weight its contribution adaptively.

Main contributions:

- A three-stream architecture integrating RGB spatial, DCT frequency, and rPPG physiological streams under a unified framework.
- The PGCMC module: an rPPG spectral-entropy gating mechanism for cross-modal consistency computation.
- CBAM-enhanced spatial stream providing both channel and spatial attention for artifact localization.
- A temporal pipeline combining frame-difference fusion, Multi-Head Self-Attention, and BiLSTM to capture local and long-range temporal cues.
- Comprehensive evaluation: ROC/PR curves, confusion matrix, noise robustness, t-SNE/PCA feature visualization, calibration, and feature importance analysis.

II. RELATED WORK

A. Spatial and Frequency Deepfake Detection

FaceForensics++ [2] established a large-scale benchmark showing that CNNs trained on RGB frames achieve high in-distribution performance but degrade under compression. Frank et al. [3] demonstrated that GAN-generated images exhibit characteristic frequency artifacts, motivating DCT-based features.




Fig. 1. PGCF architecture. Three input streams are fused by the PGCMC module; the combined representation undergoes Temporal Gradient Fusion (TGF), multi-head self-attention, BiLSTM encoding, and classification.

TABLE I
DATASET SPLIT SUMMARY

Split	Real	Fake	Total
Train	712	712	1,424
Validation	178	178	356

B. Physiological Signal-Based Detection

Li and Lyu [4] first exploited rPPG absence as a deepfake indicator. FakeCatcher [5] demonstrated real-time detection using blood volume pulse maps but remains brittle to illumination variation and motion blur.

C. Multi-Modal and Temporal Approaches

Recurrent and attention-based architectures have been applied to temporal inconsistency modeling [6]. Our work is closest to methods combining multiple evidence streams but distinguishes itself by using physiological confidence as a dynamic gating weight rather than fixed feature concatenation.

III. METHODOLOGY

A. System Overview

Given a video clip of $T = 16$ frames, PGCF extracts three parallel feature streams, fuses them through the PGCMC module, applies temporal modeling, and classifies the clip as real or fake. The enhanced temporal pipeline is:

Streams $\rightarrow$ PGCMC $\rightarrow$ TGF $\rightarrow$ Self-Attn $\rightarrow$ BiLSTM $\rightarrow$ Classifier

Figure 1 illustrates the overall architecture.

B. Spatial Stream with CBAM

The spatial stream processes $B \times T$ frames through a pretrained ResNet50 backbone. We replace the Squeeze-and-Excitation block with **CBAM** [8], applying sequential channel attention:

$$M_c(F) = \sigma(\text{MLP}(\text{AvgPool}(F)) + \text{MLP}(\text{MaxPool}(F))) \quad (1)$$

and spatial attention:

$$M_s(F') = \sigma(\text{Conv}_{7 \times 7}([\text{AvgPool}(F'); \text{MaxPool}(F')])) \quad (2)$$

where σ denotes sigmoid and $F' = M_c(F) \otimes F$. The output is projected to $\mathbf{f}_s \in \mathbb{R}^{B \times T \times 256}$.

C. Frequency Stream

A 4-channel input (RGB + DCT energy map) is processed by a 3-layer CNN ($4 \rightarrow 32 \rightarrow 64$ channels), adaptive average pooling, and an FC layer, producing $\mathbf{f}_f \in \mathbb{R}^{B \times T \times 256}$.

D. rPPG Stream

The CHROM [7] signal $\mathbf{x}_r \in \mathbb{R}^{B \times 1 \times 16}$ is encoded by a 1D-CNN ($1 \rightarrow 32 \rightarrow 64$ channels) with adaptive pooling, yielding a clip-level embedding $\mathbf{f}_r \in \mathbb{R}^{B \times 128}$.

E. PGCMC Module

rPPG Confidence. Spectral entropy H of the normalized FFT power spectrum is computed as $H = -\sum_k P_k \log(P_k + \epsilon)$; confidence is $c = \sigma(-H)$.

Cross-Modal Consistency. Pairwise cosine similarities among $\mathbf{f}_s$, $\mathbf{f}_f$, and $\mathbf{f}_r$ are averaged to produce a scalar consistency score $s \in [0, 1]$.

Gated Fusion. The concatenated feature $[\mathbf{f}_s; \mathbf{f}_f]$ is modulated by c to produce the gated feature $\tilde{\mathbf{f}} \in \mathbb{R}^{B \times T \times 512}$.

F. Temporal Modeling

Temporal Gradient Fusion (TGF) computes frame-to-frame differences and fuses them with the original features via Linear + LayerNorm + ReLU.

Temporal Self-Attention. A Transformer encoder block (4 heads, 512-dim, GELU FFN) processes the sequence:

$$\mathbf{Z} = \text{LayerNorm}(\tilde{\mathbf{f}} + \text{MHA}(\tilde{\mathbf{f}}, \tilde{\mathbf{f}}, \tilde{\mathbf{f}})) \quad (3)$$

capturing non-local temporal dependencies inaccessible to sequential models alone.

BiLSTM. The attended sequence $\mathbf{Z}$ is encoded by a bidirectional LSTM (hidden size 256), producing a 512-dim clip representation.

G. Loss Function

$$\mathcal{L} = \mathcal{L}_{\text{cls}} + 0.3 \mathcal{L}_{\text{consist}} + 0.1 \mathcal{L}_{\text{rPPG}} \quad (4)$$

$\mathcal{L}_{\text{cls}}$ is cross-entropy; $\mathcal{L}_{\text{consist}}$ penalizes low consistency for real and high consistency for fake samples; $\mathcal{L}_{\text{rPPG}}$ is a reconstruction loss on the rPPG signal.

IV. EXPERIMENTAL SETUP

A. Dataset

Each video is sampled uniformly to $T = 16$ frames at 224×224 . The validation set is balanced.

B. Training Configuration

All models are trained for 20 epochs with an effective batch size of 8 (2 per step, 4 gradient accumulation steps). Two learning rate groups are used: 10^{-4} for the pretrained backbone and 5×10^{-4} for new modules. AdamW with weight decay 10^{-4} and CosineAnnealingLR ($\eta_{\min} = 10^{-6}$) are used. Mixed-precision training (`torch.cuda.amp`) is enabled. The ResNet50 backbone is progressively unfrozen at epoch 3; auxiliary losses are warmed up from epoch 3.

TABLE II
PGCF EVALUATION METRICS ON BALANCED VALIDATION SET

Metric	Value
Accuracy	0.7416
AUC (ROC)	0.8240
F1 Score	0.7416
Precision	0.7416
Recall	0.7416
Avg. Precision (PR)	0.7907



Fig. 2. ROC Curve for the PGCF DeepFake Detector. The model achieves an AUC of 0.8240, demonstrating strong discriminative capability over the random baseline.

V. RESULTS

A. Classification Performance

Table II reports evaluation metrics on the balanced 356-sample validation set.

The per-class breakdown confirms balanced classification: both Real and Fake classes achieve a precision, recall, and F1 of 0.74, with 132 correct predictions out of 178 samples per class. The AUC of 0.8240 indicates strong ranking ability well above the random baseline.

B. ROC and Precision-Recall Curves

The ROC curve (Fig. 2) shows the model reaching approximately 0.60 TPR at 0.20 FPR, a region relevant for practical deployment with limited false positive tolerance. The AUC of 0.8240 confirms strong discriminability. The PR curve (Fig. 3) maintains precision above 0.80 for recall up to ~ 0.40 , then gradually degrades, yielding AP = 0.7907.

C. Confusion Matrix

The confusion matrix shows symmetric classification behaviour:

$$C = \begin{pmatrix} 132 & 46 \\ 46 & 132 \end{pmatrix}$$

yielding a 74.2% true positive rate for both Real and Fake classes. The symmetric error pattern indicates no class bias despite the balanced training set.

3D PCA — PGCF BiLSTM Features




Fig. 3. Precision-Recall Curve for PGCF on the balanced validation set. The detector achieves an Average Precision (AP) of 0.7907 while maintaining high precision across moderate recall regions.




Fig. 4. Confusion matrix of the PGCF detector on the balanced validation set. Both Real and Fake classes achieve symmetric classification performance with 132 correct predictions each.

D. Cross-Modal Consistency Analysis

Table III summarizes the PGCMC consistency score statistics on the validation set. The gap of 0.418 between class means validates the central hypothesis: genuine faces maintain higher cross-modal alignment across spatial, frequency, and physiological streams.

E. Training Dynamics

The PGCMC consistency scores converge monotonically: real videos stabilize near 0.83 by epoch 20, while fake videos drop to approximately 0.15. Training accuracy reaches 91.5% by epoch 20 while validation accuracy plateaus at $\approx 74\%$,




Fig. 5. Distribution of PGCMC cross-modal consistency scores for Real and Fake videos. Genuine videos exhibit substantially higher consistency values, validating the physiological coherence hypothesis underlying PGCF.

TABLE III
PGCMC CONSISTENCY SCORE STATISTICS

Class	Mean (μ)	Ideal
Real	0.690	$\rightarrow 1.0$
Fake	0.272	$\rightarrow 0.0$
Gap	0.418	

TABLE IV
TRAINING PROGRESSION HIGHLIGHTS

Epoch	Train Acc	Val Acc	Val AUC	Val F1
3	0.545	0.685	0.732	0.705
6	0.736	0.708	0.800	0.741
9	0.801	0.773	0.824	0.777
14	0.872	0.773	0.843	0.771
17	0.867	0.770	0.846	0.771
20	0.915	0.767	0.842	0.786

indicative of overfitting on the limited training set of 1,424 clips.

Table IV reports key training milestones.

F. Robustness Evaluation

Table V reports performance under additive Gaussian noise. The model exhibits graceful degradation: at $\sigma = 0.10$, AUC remains at 0.77; at $\sigma = 0.20$, AUC is 0.73. This robustness is attributed to the physiological gating in PGCMC, which anchors consistency even when the spatial stream is partially corrupted.

G. Feature Space Analysis

t-SNE visualization of BiLSTM output features shows partial class separation: real videos occupy a curved manifold while fake videos cluster more diffusely, with overlap in the central region corresponding to hard-to-classify samples. **3D PCA** reveals PC1 explaining 63.7% and PC2 explaining 10.2%

TABLE V
ROBUSTNESS TO GAUSSIAN INPUT NOISE

Noise σ	AUC	Accuracy
0.000	0.82	0.74
0.010	0.82	0.74
0.025	0.83	0.73
0.050	0.79	0.70
0.100	0.77	0.66
0.150	0.74	0.65
0.200	0.73	0.62

of variance, with two distinct linear subspaces along PC1 consistent with a primary discriminant axis aligned to the real-fake boundary.

H. Feature Importance

A Random Forest fitted to frozen BiLSTM feature vectors identifies Feature 68 as the most discriminative (importance ≈ 0.025), substantially higher than Features 165, 19, and 58 (all ≈ 0.013 – 0.015). The remaining top-30 features carry lower but non-trivial importance, indicating distributed rather than concentrated representation.

I. Calibration Analysis

The calibration curve reveals moderate miscalibration, particularly at low (0.05–0.20) and mid (0.30–0.50) predicted probability ranges. Above 0.60, the model approaches the diagonal. Post-hoc temperature scaling or Platt calibration is recommended for deployment scenarios requiring reliable confidence estimates.

VI. DISCUSSION

A. Strengths

The PGCMC consistency gap (0.418 between class means) validates the physiological grounding of PGCF. The AUC of 0.8240 demonstrates strong discriminability on a balanced evaluation set. Robustness to Gaussian noise up to $\sigma = 0.20$ without catastrophic collapse is a positive indicator for real-world deployment.

B. Limitations and Future Work

Overfitting. The 17.3 percentage-point gap between training (91.5%) and validation (74.2%) accuracy indicates that stronger regularization—mixup augmentation, increased dropout, stochastic depth—should be explored.

Dataset scale. With 1,424 training clips, cross-dataset generalization on FaceForensics++ [2] or DFDC remains future work.

Calibration. Moderate miscalibration at low-confidence bins warrants post-hoc calibration for threshold-sensitive applications.

Component ablation. Isolating the contribution of CBAM vs. temporal self-attention through controlled ablation is planned as the PGCF v2 enhanced model awaits execution.




Fig. 6. Training dynamics of PGCF across 20 epochs. The figure shows loss convergence, accuracy evolution, validation AUC/F1 progression, and PGCMC consistency score separation between Real and Fake classes.

VII. CONCLUSION

We presented PGCF, a physiology-guided three-stream deepfake detection framework fusing spatial, frequency, and rPPG signals through a cross-modal consistency module. The PGCMC gating mechanism demonstrates that physiological coherence is a measurable and effective discriminant, with a 0.418 gap in mean consistency scores between real and fake videos. On a balanced validation set of 356 clips, PGCF achieves an AUC of 0.8240, accuracy of 74.16%, and Average Precision of 0.7907, while retaining an AUC of 0.73 under heavy Gaussian noise. Future work will focus on cross-dataset generalization, ablation of architectural components, and post-hoc calibration for deployment readiness.

REFERENCES

- [1] R. Tolosana, R. Vera-Rodriguez, J. Fierrez, A. Morales, and J. Ortega-Garcia, "Deepfakes and beyond: A survey of face manipulation and fake detection," *Inf. Fusion*, vol. 64, pp. 131–148, 2020.
- [2] A. Rössler, D. Cozzolino, L. Verdoliva, C. Riess, J. Thies, and M. Nießner, "FaceForensics++: Learning to detect manipulated facial images," in *Proc. IEEE/CVF ICCV*, 2019, pp. 1–11.
- [3] J. Frank, T. Eisenhofer, L. Schönherr, A. Fischer, D. Kolossa, and T. Holz, "Leveraging frequency analysis for deep fake image recognition," in *Proc. ICML*, 2020.
- [4] X. Li and S. Lyu, "Exposing DeepFake videos by detecting face warping artifacts," in *Proc. IEEE CVPR Workshops*, 2019.
- [5] U. A. Ciftci, I. Demir, and L. Yin, "FakeCatcher: Detection of synthetic portrait videos using biological signals," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 44, no. 9, pp. 4458–4470, 2020.
- [6] E. Sabir, J. Cheng, A. Jaiswal, W. AbdAlmageed, I. Masi, and P. Natarajan, "Recurrent convolutional strategies for face manipulation detection in videos," *Interfaces (GUI)*, vol. 3, no. 1, pp. 80–87, 2019.
- [7] G. de Haan and V. Jeanne, "Robust pulse rate from chrominance-based rPPG," *IEEE Trans. Biomed. Eng.*, vol. 60, no. 10, pp. 2878–2886, 2013.
- [8] S. Woo, J. Park, J.-Y. Lee, and I. S. Kweon, "CBAM: Convolutional block attention module," in *Proc. ECCV*, 2018, pp. 3–19.




Fig. 7. Robustness analysis of PGCF under additive Gaussian noise. The model exhibits graceful degradation in both AUC and accuracy as noise variance increases.




Fig. 8. t-SNE visualization of BiLSTM feature embeddings generated by PGCF. Real and Fake samples form partially separable manifolds, with overlap concentrated near hard-to-classify samples.




Fig. 9. 3D PCA projection of PGCF BiLSTM features. PC1 explains 63.7% of the variance and reveals two dominant subspaces corresponding to real and fake representations.




Fig. 10. Top-30 Random Forest feature importances computed from frozen BiLSTM embeddings. Feature 68 emerges as the most discriminative component for deepfake classification.




Fig. 11. Calibration curve for PGCF. The detector approaches good calibration at higher confidence ranges, while low-confidence regions exhibit moderate probability mismatch.